

DQN_v1

August 15, 2020

0.1 Jason's toy model in Q-learning

Thanks for this useful link:https://www.tensorflow.org/agents/tutorials/1_dqn_tutorial So far, Reinforcement learning algorithms are not well GPU optimized since it's hard to do parallization in updating stages + reward. Thus, the DQN part is CPU only :) Thanks to this useful link : <https://adventuresinmachinelearning.com/reinforcement-learning-tensorflow/> Here we didn't consider the Volume

```
In [1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib
import tensorflow as tf
from matplotlib import colors as mcolors
import matplotlib.pyplot as plt
color_array = list(mcolors.CSS4_COLORS.keys())
root_path = 'Data/Stocks/'
#root_path = "/kaggle/input/price-volume-data-for-all-us-stocks-etfs/ETFs/"

## Here we got 2 sets of data: Google stock price daily and Google stock price in 5 da

df = pd.read_csv(root_path + "googl.us.txt")
df_hf = pd.read_csv("Data/High-frequency_data/GOOG_stock_1minute_sample/GOOG_sample.txt")
```

```
In [2]: df.head()
```

```
Out[2]:
```

	Date	Open	High	Low	Close	Volume	OpenInt
0	2004-08-19	50.000	52.03	47.980	50.170	44703800	0
1	2004-08-20	50.505	54.54	50.250	54.155	22857200	0
2	2004-08-23	55.375	56.74	54.525	54.700	18274400	0
3	2004-08-24	55.620	55.80	51.785	52.435	15262600	0
4	2004-08-25	52.480	54.00	51.940	53.000	9197800	0

```
In [3]: df_hf.head()
```

```
Out[3]:
```

	DateTime	Open	High	Low	Close	Volume
0	2020-01-02 04:00:00	1342.00	1342.20	1342.00	1342.20	424
1	2020-01-02 04:02:00	1344.20	1344.20	1344.20	1344.20	177
2	2020-01-02 08:00:00	1337.02	1337.02	1337.02	1337.02	329
3	2020-01-02 08:09:00	1347.00	1347.00	1347.00	1347.00	155
4	2020-01-02 08:55:00	1348.00	1348.00	1348.00	1348.00	190

```

In [4]: import matplotlib
        from matplotlib.pylab import rc
        font = {'family': 'normal', 'weight': 'bold',
                'size': 25}

        matplotlib.rc('font', **font)
        rc('axes', linewidth=3)

        plt.subplot(1,1,1)
        plt.plot(df["Open"], "k", label="Open", alpha=0.9)
        plt.plot(df["Close"], "r", label="Close", alpha=0.9)


        plt.xlabel("Day")
        plt.ylabel(r"${\rm value}$")
        plt.suptitle("Stock price vs Day Google")

        fig = matplotlib.pyplot.gcf()

        fig.set_size_inches(20,10)
        plt.legend(fontsize=25, handlelength=5, ncol=3)
        plt.show()

```

findfont: Font family ['normal'] not found. Falling back to DejaVu Sans.
 findfont: Font family ['normal'] not found. Falling back to DejaVu Sans.
 findfont: Font family ['normal'] not found. Falling back to DejaVu Sans.



```

In [5]: import matplotlib
        from matplotlib.pylab import rc
        font = {'family': 'normal', 'weight': 'bold',
                'size': 25}

        matplotlib.rc('font', **font)
        rc('axes', linewidth=3)


        plt.plot(df_hf["Open"], "k", label= "Open")
        plt.plot(df_hf["Close"], "r", label= "Close")

        plt.xlabel("Minutes")
        plt.ylabel(r"${\rm value}$")
        plt.suptitle("Stock price vs Minute Google including close market")
        plt.legend()

        fig = matplotlib.pyplot.gcf()
        fig.set_size_inches(20,10)

```

Stock price vs Minute Google including close market



```

In [6]: df_hf["DateTime"] = pd.to_datetime(df_hf["DateTime"])
        delta_min = (df_hf["DateTime"]-df_hf["DateTime"][0]).dt.total_seconds()/60

In [7]: import matplotlib
        from matplotlib.pylab import rc
        font = {'family': 'normal', 'weight': 'bold',
                'size': 25}

```

```

matplotlib.rc('font', **font)
rc('axes', linewidth=3)


plt.plot(delta_min,df_hf["Open"],"kx",label= "Open")
plt.plot(delta_min,df_hf["Close"],"ro",label= "Close")

plt.xlabel("Minutes")
plt.ylabel(r"${\rm value}$")
plt.suptitle("Stock price vs Minute Google")
plt.legend()

fig = matplotlib.pyplot.gcf()
fig.set_size_inches(20,10)

```

Stock price vs Minute Google



In [8]: *# This time Try DQN on df and df_hf*

```

class Environment1:

    def __init__(self, data, history_t=90):
        self.data = data
        self.history_t = history_t
        self.reset()

    def reset(self):
        self.t = 0
        self.done = False
        self.profits = 0

```

```

self.positions = []
self.position_value = 0
self.history = [0 for _ in range(self.history_t)]
return [self.position_value] + self.history # obs

def step(self, act):
    reward = 0

    # act = 0: stay, 1: buy, 2: sell
    if act == 1:
        self.positions.append(self.data.iloc[self.t, :]['Close'])
    elif act == 2: # sell
        if len(self.positions) == 0:
            reward = -1
        else:
            profits = 0
            for p in self.positions:
                profits += (self.data.iloc[self.t, :]['Close'] - p)
            reward += profits
            self.profits += profits
            self.positions = []

    # set next time
    self.t += 1
    self.position_value = 0
    for p in self.positions:
        self.position_value += (self.data.iloc[self.t, :]['Close'] - p)
    self.history.pop(0)
    self.history.append(self.data.iloc[self.t, :]['Close'] - self.data.iloc[(self.t-1), :]['Close'])

    # clipping reward
    if reward > 0:
        reward = 1
    elif reward < 0:
        reward = -1

    return [self.position_value] + self.history, reward, self.done

```

In [9]: # df

```

df['Date'] = pd.to_datetime(df['Date'])
data = df.set_index('Date')
env = Environment1(df)

```

```

# df_hf

```

```
data_hf = df_hf.set_index('DateTime')
env_hf = Environment1(df_hf)
```

In [10]: *# Train test split :) Here I didn't shuffle dataset since we want to check whether we*

```
delta_split = int(df.shape[0]*0.7)

train = df[:delta_split]
test = df[delta_split:]
print("Training date length",len(train),"testing date length", len(test))
```

Training date length 2333 testing date length 1000

```
In [11]: delta_split = int(df_hf.shape[0]*0.7)
train_hf = df_hf[:delta_split]
test_hf = df_hf[delta_split:]
print("Training date length hf",len(train_hf),"testing date length hf", len(test_hf))
```

Training date length hf 2275 testing date length hf 976

```
In [13]: import chainer
import chainer.functions as F
import chainer.links as L
import copy
import time

def train_dqn(env):

    class Q_Network(chainer.Chain):

        def __init__(self, input_size, hidden_size, output_size):
            super(Q_Network, self).__init__(
                fc1 = L.Linear(input_size, hidden_size),
                fc2 = L.Linear(hidden_size, hidden_size),
                fc3 = L.Linear(hidden_size, output_size)
            )

        def __call__(self, x):
            h = F.relu(self.fc1(x))
            h = F.relu(self.fc2(h))
            y = self.fc3(h)
            return y

        def reset(self):
```

```

        self.zerograds()

Q = Q_Network(input_size=env.history_t+1, hidden_size=100, output_size=3)
# GPU Option
#Q.to_gpu()

Q_ast = copy.deepcopy(Q)
optimizer = chainer.optimizers.Adam()
optimizer.setup(Q)

# Hyper-parameters
epoch_num = 25
step_max = len(env.data)-1
memory_size = 200
batch_size = 20
epsilon = 1.0
epsilon_decrease = 1e-3
epsilon_min = 0.1
start_reduce_epsilon = 200
train_freq = 10
update_q_freq = 20
gamma = 0.97
show_log_freq = 5

memory = []
total_step = 0
total_rewards = []
total_losses = []

start = time.time()
for epoch in range(epoch_num):

    pobs = env.reset()
    step = 0
    done = False
    total_reward = 0
    total_loss = 0

    while not done and step < step_max:

        # select act
        pact = np.random.randint(3)
        if np.random.rand() > epsilon:
            pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
            pact = np.argmax(pact.data)

        # act

```

```

obs, reward, done = env.step(pact)

# add memory
memory.append((pobs, pact, reward, obs, done))
if len(memory) > memory_size:
    memory.pop(0)

# train or update q
if len(memory) == memory_size:
    if total_step % train_freq == 0:
        shuffled_memory = np.random.permutation(memory)
        memory_idx = range(len(shuffled_memory))
        for i in memory_idx[:batch_size]:
            batch = np.array(shuffled_memory[i:i+batch_size])
            b_pobs = np.array(batch[:, 0].tolist(), dtype=np.float32).reshape(batch_size, 1)
            b_pact = np.array(batch[:, 1].tolist(), dtype=np.int32).reshape(batch_size, 1)
            b_reward = np.array(batch[:, 2].tolist(), dtype=np.int32).reshape(batch_size, 1)
            b_obs = np.array(batch[:, 3].tolist(), dtype=np.float32).reshape(batch_size, 1)
            b_done = np.array(batch[:, 4].tolist(), dtype=np.bool).reshape(batch_size, 1)

            q = Q(b_pobs)
            maxq = np.max(Q_ast(b_obs).data, axis=1)
            target = copy.deepcopy(q.data)
            for j in range(batch_size):
                target[j, b_pact[j]] = b_reward[j] + gamma * maxq[j] * (not b_done[j])
            Q.reset()
            loss = F.mean_squared_error(q, target)
            total_loss += loss.data
            loss.backward()
            optimizer.update()

        if total_step % update_q_freq == 0:
            Q_ast = copy.deepcopy(Q)

# epsilon
if epsilon > epsilon_min and total_step > start_reduce_epsilon:
    epsilon -= epsilon_decrease

# next step
total_reward += reward
pobs = obs
step += 1
total_step += 1

total_rewards.append(total_reward)
total_losses.append(total_loss)

if (epoch+1) % show_log_freq == 0:

```



```

log_reward = sum(total_rewards[((epoch+1)-show_log_freq):])/show_log_freq
log_loss = sum(total_losses[((epoch+1)-show_log_freq):])/show_log_freq
elapsed_time = time.time()-start
print('\t'.join(map(str, [epoch+1, epsilon, total_step, log_reward, log_loss])))
start = time.time()

```

```

return Q, total_losses, total_rewards

```

In [14]: # Train on google_daily:

```

# Q, total_losses, total_rewards = train_dqn(Environment1(train))

```

In [15]: # Train on google_per_minute:

```

# Epoch, epsilon (Randomness in your strategy), steps, log[reward], log[loss], elapsed_time
Q, total_losses, total_rewards = train_dqn(Environment1(train_hf))

```

5	0.09999999999999992	11370	-15.4	5297.414914591704	113.0948
10	0.09999999999999992	22740	23.4	121.3875683060789	111.3312
15	0.09999999999999992	34110	37.6	169.700724235503	114.35927
20	0.09999999999999992	45480	24.2	62.293562510469926	111.294
25	0.09999999999999992	56850	35.0	44.13690594714135	113.9022

```

In [20]: plt.subplot(1,2,1)
plt.plot(total_losses,"r",label="loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()

plt.subplot(1,2,2)
plt.plot(total_rewards,"b",label="Reward")
plt.xlabel("Epoch")
plt.ylabel("Reward")

plt.legend()

```

```

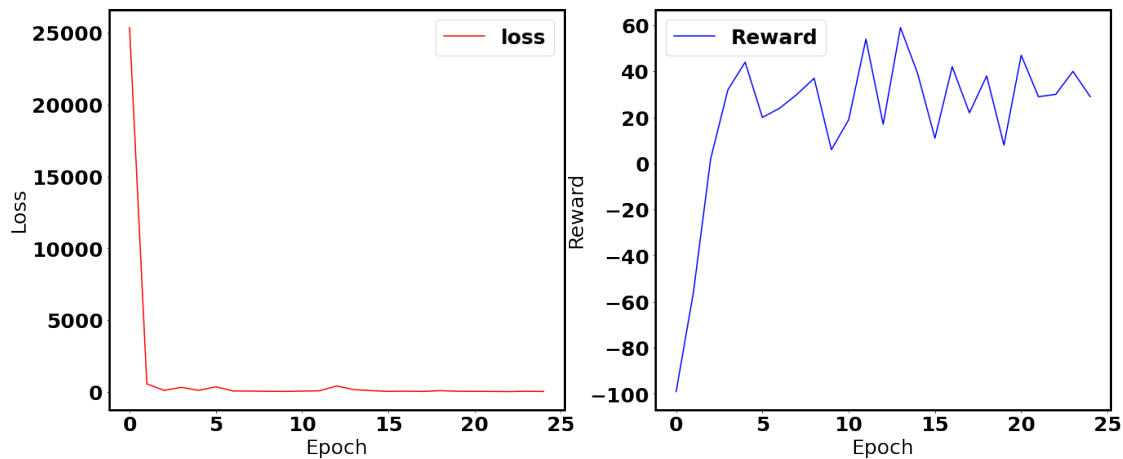
fig = matplotlib.pyplot.gcf()

```

```

fig.set_size_inches(22,9)
plt.legend()
plt.show()

```



In [40]: # Testing :)

```
# test
test_env = Environment1(test_hf)

# test
pobs = test_env.reset()
test_acts = []
test_rewards = []

for _ in range(len(test_env.data) - 1):
    pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
    pact = np.argmax(pact.data)
    test_acts.append(pact)

    obs, reward, done = test_env.step(pact)
    test_rewards.append(reward)

    pobs = obs

test_profits = test_env.profits

print("Test profit =%.4f in %d steps start price %.4f"%(test_profits, test_env.data.shape[0], test_env.data[0]))

Test profit =117.8347 in 976 steps start price 1416.0400


In [69]: # Our actions in testing set:
# Here 0 means stay, 1 means buy 2 means sell
plt.hist(test_acts, label = "0 stay 1 buy 2 sell")
```

```

fig = matplotlib.pyplot.gcf()

fig.set_size_inches(22,10)
plt.legend()
plt.show()

```



```

In [68]: plt.subplot(1,1,1)
         y = test_env.data["Close"].values[1:]

         test_acts = np.array(test_acts)


         m0 = test_acts==0
         m1 = test_acts==1
         m2 = test_acts==2
         plt.plot(np.arange(0,len(y),1),y,label = "Close price")
         plt.plot(np.arange(0,len(y),1)[m0],y[m0], "ko",label = "Stay")
         plt.plot(np.arange(0,len(y),1)[m1],y[m1], "ro",label = "Buy")
         plt.plot(np.arange(0,len(y),1)[m2],y[m2], "go",label = "Sell")

         plt.xlabel("Minutes ")
         plt.ylabel("Price")

         fig = matplotlib.pyplot.gcf()

         fig.set_size_inches(22,10)
         plt.legend()
         plt.show()

```



In []:

In []:

In []:

In []: