

Stock_Qlearning

August 3, 2020

0.1 Deep learning/Reinforcement learning in Time series data: Stock price verion

Our goal is to predict close price using our model from previous input Episode one: <https://www.kaggle.com/peraktong/transformer-ts> Please let me know if you have any question. Thanks to this useful link : <https://www.tensorflow.org/tutorials/text/transformer> Thanks to this useful link: <https://www.kaggle.com/itoeiji/deep-reinforcement-learning-on-stock-data> Thanks to this useful link: <https://towardsdatascience.com/aifortrading-2edd6fac689d> Thanks to this useful link: https://www.tensorflow.org/agents/tutorials/1_dqn_tutorial Thanks to this useful link: <https://rubikscore.net/2019/07/08/deep-q-learning-with-python-and-tensorflow-2-0/>

```
In [1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib
import tensorflow as tf
from matplotlib import colors as mcolors
import matplotlib.pyplot as plt
color_array = list(mcolors.CSS4_COLORS.keys())
# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all fi

import os
names_list = []
for dirname, _, filenames in os.walk('Data/Stocks/'):
    for filename in filenames:
        #print(os.path.join(dirname, filename))

        names_list.append(filename)
names_list = np.array(names_list)

# TPU session :)
# detect and init the TPU
if False:
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver()
    tf.config.experimental_connect_to_cluster(tpu)
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
# instantiate a distribution strategy
tpu_strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
In [2]: root_path = 'Data/Stocks/'
        #root_path = "/kaggle/input/price-volume-data-for-all-us-stocks-etfs/ETFs/"

        df = pd.read_csv(root_path + "googl.us.txt")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Date	Open	High	Low	Close	Volume	OpenInt
0	2004-08-19	50.000	52.03	47.980	50.170	44703800	0
1	2004-08-20	50.505	54.54	50.250	54.155	22857200	0
2	2004-08-23	55.375	56.74	54.525	54.700	18274400	0
3	2004-08-24	55.620	55.80	51.785	52.435	15262600	0
4	2004-08-25	52.480	54.00	51.940	53.000	9197800	0

```
In [4]: import matplotlib
        from matplotlib.pylab import rc
        font = {'family': 'normal', 'weight': 'bold',
                'size': 25}

        matplotlib.rc('font', **font)
        rc('axes', linewidth=3)

        plt.subplot(1,1,1)
        plt.plot(df["Open"], "k", label="Open")
        plt.plot(df["Close"], "r", label="Close")

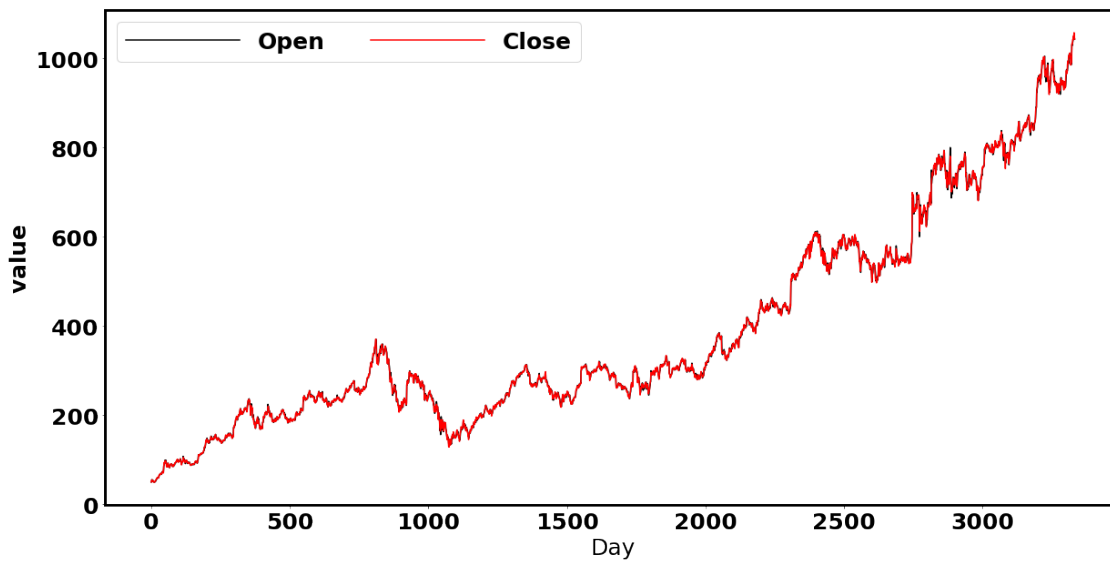
        plt.xlabel("Day")
        plt.ylabel(r"${\rm value}$")
        plt.suptitle("Stock price vs Day Google")

        fig = matplotlib.pyplot.gcf()

        fig.set_size_inches(20,10)
        plt.legend(fontsize=25, handlelength=5, ncol=3)
        plt.show()
```

```
findfont: Font family ['normal'] not found. Falling back to DejaVu Sans.
findfont: Font family ['normal'] not found. Falling back to DejaVu Sans.
findfont: Font family ['normal'] not found. Falling back to DejaVu Sans.
```

Stock price vs Day Google



```
In [5]: # Let's check whether there is a peroid in week/month?
import datetime
df['Date'] = pd.to_datetime(df['Date'], errors='coerce')
df['weekday'] = df['Date'].dt.weekday
df['monthday'] = df['Date'].dt.day
df['month'] = df['Date'].dt.month

from matplotlib.pylab import rc
font = {'family': 'normal', 'weight': 'bold',
        'size': 25}

matplotlib.rc('font', **font)
rc('axes', linewidth=3)

color_array = ['b', 'g', 'r', 'c', 'm', 'y', 'k']

plt.subplot(1,1,1)

names_array = list(df.keys()[1:5])
for i in range(len(names_array)):
    #print("Doing %d"%i)
    plt.plot(df.groupby('weekday').mean().index+1, df.groupby('weekday').mean()[names_array[i]])
```

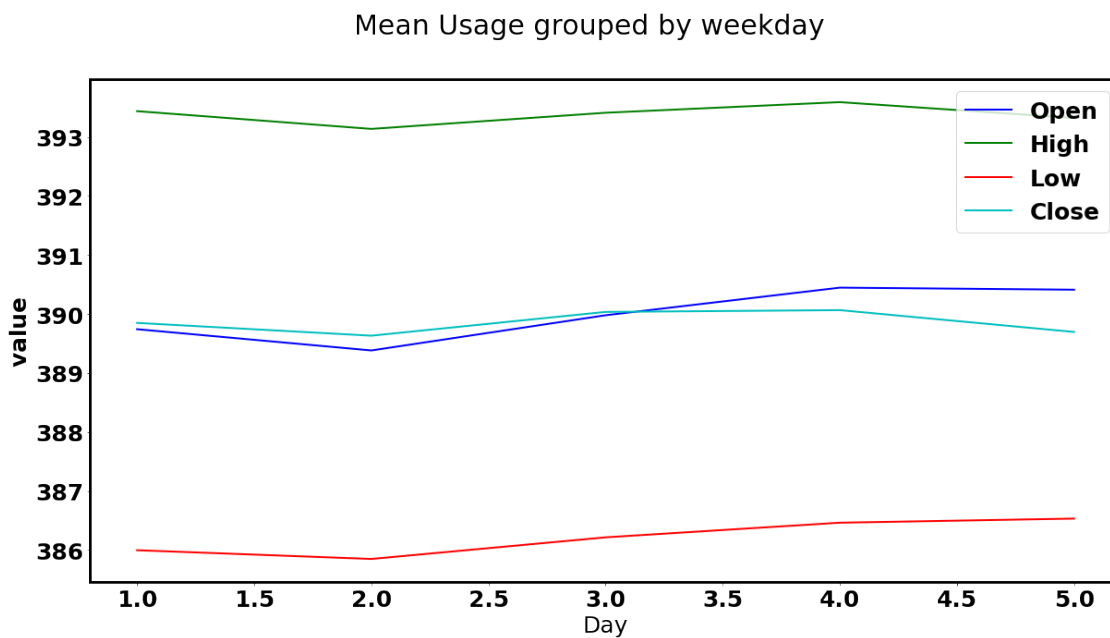
```

plt.xlabel("Day")
plt.ylabel(r"${\rm value}$")
plt.suptitle("Mean Usage grouped by weekday")

fig = matplotlib.pyplot.gcf()

fig.set_size_inches(20,10)
plt.legend()
plt.show()

```



```

In [6]: from matplotlib.pylab import rc
        font = {'family': 'normal', 'weight': 'bold',
                'size': 25}

        matplotlib.rc('font', **font)
        rc('axes', linewidth=3)

        color_array = ['b', 'g', 'r', 'c', 'm', 'y', 'k']

        plt.subplot(1,1,1)

        names_array = list(df.keys()[1:5])
        for i in range(len(names_array)):
            #print("Doing %d"%i)
            plt.plot(df.groupby('monthday').mean().index+1, df.groupby('monthday').mean()[names.

```


```

plt.xlabel("Day")
plt.ylabel(r"${\rm value}$")
plt.suptitle("Mean Usage grouped by monthday")

fig = matplotlib.pyplot.gcf()

fig.set_size_inches(20,10)
plt.legend()
plt.show()

```



In [7]: # By month?

```

from matplotlib.pylab import rc
font = {'family': 'normal', 'weight': 'bold',
        'size': 25}

matplotlib.rc('font', **font)
rc('axes', linewidth=3)

color_array = ['b', 'g', 'r', 'c', 'm', 'y', 'k']

```

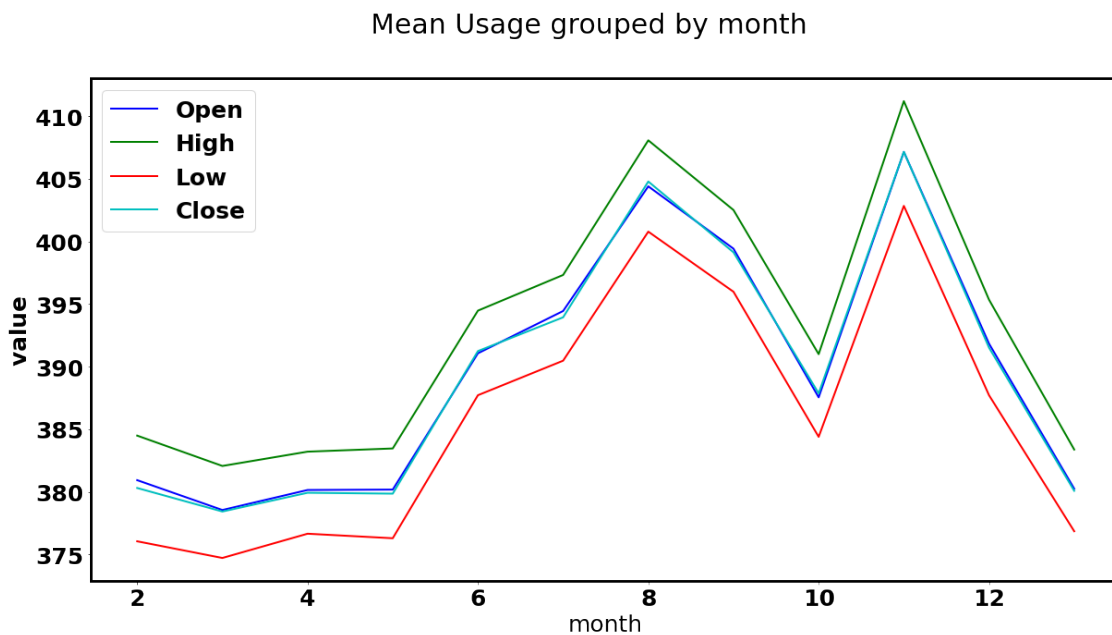
```
plt.subplot(1,1,1)
```

```
names_array = list(df.keys()[1:5])
for i in range(len(names_array)):
    #print("Doing %d"%i)
    plt.plot(df.groupby('month').mean().index+1,df.groupby('month').mean()[names_array
```

```
plt.xlabel("month")
plt.ylabel(r"${\rm value}$")
plt.suptitle("Mean Usage grouped by month")
```

```
fig = matplotlib.pyplot.gcf()
```

```
fig.set_size_inches(20,10)
plt.legend()
plt.show()
```



1 baseline model: GRU+Dense layer structure:

```
In [12]: from sklearn import preprocessing
         from keras.models import Model
```

```

from sklearn.preprocessing import LabelEncoder
import tensorflow as tf
from keras.layers import Input,Masking,Dense,GRU
from sklearn.model_selection import train_test_split
from keras.layers import Dropout

In [13]: # some hyper parameters
# Here we use 30 day data to predict future stock price. May need to modify later
delta_t = 30
n_cell = 50
dropout_rate= 0.1
n_epoch=30
batch=256

In [14]: checkpoint_path = "model/cp.ckpt"
checkpoint_dir = os.path.dirname(checkpoint_path)

min_max_scaler = preprocessing.StandardScaler()

# min-max scaler
np_scaled = min_max_scaler.fit_transform(df[names_array])

df_scaled = pd.DataFrame(np_scaled,columns=names_array)

X = np.zeros((df_scaled.shape[0]-delta_t,delta_t,len(names_array)),dtype=float)
y = df_scaled["Close"][delta_t:]

for i in range(len(y)):
    if i%800==0:
        print("Prepare data %.2f percent"%(100*i/len(y)))
        X[i,.,:] = df_scaled[i:i+delta_t][names_array].values

# split train test:
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, shuffle=True)

Prepare data 0.00 percent
Prepare data 24.22 percent
Prepare data 48.44 percent
Prepare data 72.66 percent
Prepare data 96.88 percent

In [15]: def generate_model(MAX_TIMESTEPS,MAX_NB_VARIABLES):
    ip = Input(shape=(MAX_TIMESTEPS,MAX_NB_VARIABLES))
    # split into x and y two channels

    x = Masking()(ip)

```

```

x = GRU(n_cell)(x)

out = Dense(1)(x)
#out = Dropout(dropout_rate)(x)
#loc, scale = tf.keras.layers.GaussianNoise(stddev=0.01)(out)

# now we only output loc
#out=loc

model = Model(ip, out)
model.summary()

# add load model code here to fine-tune

return model

model = generate_model(delta_t,X.shape[2])
model.compile(loss='mae', optimizer='adam')
#model.summary()

callback = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path,
                                              save_weights_only=True,
                                              verbose=1)

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 30, 4)]	0
masking (Masking)	(None, 30, 4)	0
gru (GRU)	(None, 50)	8400
dense (Dense)	(None, 1)	51
Total params: 8,451		
Trainable params: 8,451		
Non-trainable params: 0		

In [16]: history = model.fit(X_train, y_train, epochs=n_epoch, batch_size=batch, validation_data=(X_val, y_val))

Epoch 1/30

Epoch 00001: saving model to model/cp.ckpt
10/10 - 2s - loss: 0.6025 - val_loss: 0.2666

Epoch 2/30

Epoch 00002: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.1754 - val_loss: 0.1587
Epoch 3/30

Epoch 00003: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.1070 - val_loss: 0.0618
Epoch 4/30

Epoch 00004: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0652 - val_loss: 0.0540
Epoch 5/30

Epoch 00005: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0576 - val_loss: 0.0427
Epoch 6/30

Epoch 00006: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0432 - val_loss: 0.0430
Epoch 7/30

Epoch 00007: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0424 - val_loss: 0.0387
Epoch 8/30

Epoch 00008: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0363 - val_loss: 0.0361
Epoch 9/30

Epoch 00009: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0341 - val_loss: 0.0319
Epoch 10/30

Epoch 00010: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0315 - val_loss: 0.0302
Epoch 11/30

Epoch 00011: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0292 - val_loss: 0.0283
Epoch 12/30

Epoch 00012: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0275 - val_loss: 0.0268
Epoch 13/30

Epoch 00013: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0261 - val_loss: 0.0257

Epoch 14/30

Epoch 00014: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0256 - val_loss: 0.0252
Epoch 15/30

Epoch 00015: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0257 - val_loss: 0.0255
Epoch 16/30

Epoch 00016: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0266 - val_loss: 0.0257
Epoch 17/30

Epoch 00017: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0270 - val_loss: 0.0249
Epoch 18/30

Epoch 00018: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0269 - val_loss: 0.0241
Epoch 19/30

Epoch 00019: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0261 - val_loss: 0.0240
Epoch 20/30

Epoch 00020: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0253 - val_loss: 0.0239
Epoch 21/30

Epoch 00021: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0250 - val_loss: 0.0240
Epoch 22/30

Epoch 00022: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0255 - val_loss: 0.0238
Epoch 23/30

Epoch 00023: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0257 - val_loss: 0.0236
Epoch 24/30

Epoch 00024: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0256 - val_loss: 0.0235
Epoch 25/30

Epoch 00025: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0252 - val_loss: 0.0236

Epoch 26/30

Epoch 00026: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0249 - val_loss: 0.0233
Epoch 27/30

Epoch 00027: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0251 - val_loss: 0.0233
Epoch 28/30

Epoch 00028: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0249 - val_loss: 0.0232
Epoch 29/30

Epoch 00029: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0248 - val_loss: 0.0233
Epoch 30/30

Epoch 00030: saving model to model/cp.ckpt
10/10 - 0s - loss: 0.0245 - val_loss: 0.0231

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, shuffle=False)
        y_pre = model.predict(X_test)
```

```
In [19]: from matplotlib.pylab import rc
        font = {'family': 'normal', 'weight': 'bold',
                'size': 25}

        matplotlib.rc('font', **font)
        rc('axes', linewidth=3)

        color_array = ['b', 'g', 'r', 'c', 'm', 'y', 'k']

        plt.subplot(1,1,1)

        min_val = np.nanmin(df["Close"])
        max_val = np.nanmax(df["Close"])

        date_array = np.arange(0, len(y_test), 1)

        y1 = min_val + (max_val - min_val) * y_test
        y2 = min_val + (max_val - min_val) * y_pre[:, 0]
        plt.plot(date_array, y1, "k", label = "Data")
        plt.plot(date_array, y2, "r", label = "Prediction")
        plt.plot(date_array, y1 - y2, "b", label = "Difference")
```

```

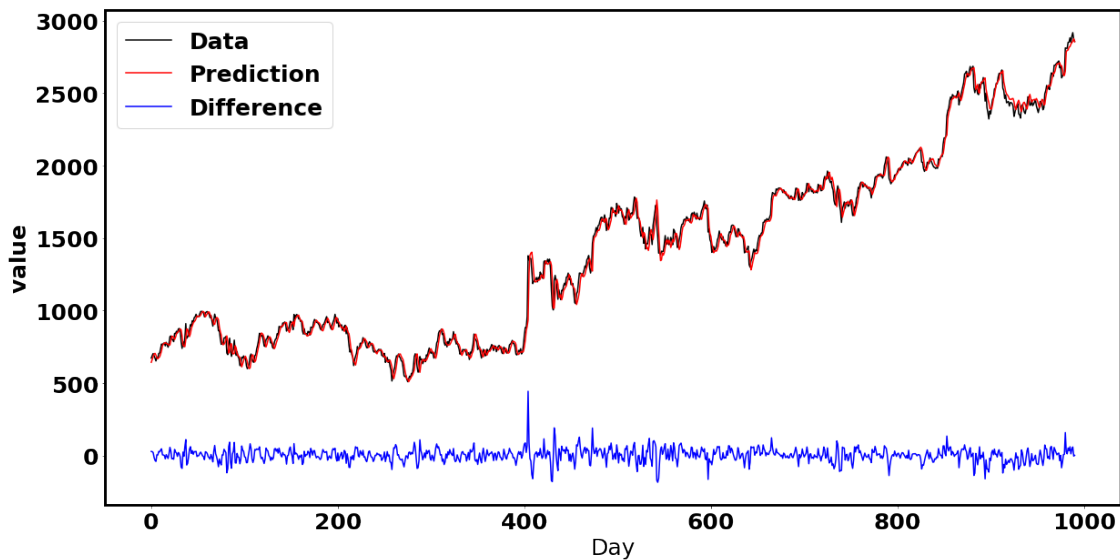
plt.xlabel("Day")
plt.ylabel(r"${\rm value}$")
plt.suptitle("Stock price prediction MAE=%.4f"%np.nanmean(abs(y1-y2)))

fig = matplotlib.pyplot.gcf()

fig.set_size_inches(20,10)
plt.legend()
plt.show()

```

Stock price prediction MAE=33.9523



1.1 The prediction seems to have a "Latency" inside.

1.2 Let's try the Q-net

In Q-learning we learn the value of taking an action from a given state. Q-value is the expected return after taking the action. We will use Rainbow which is a combination of seven Q learning algorithms.

In [8]: `class Environment1:`

```

def __init__(self, data, history_t=90):
    self.data = data
    self.history_t = history_t
    self.reset()

def reset(self):
    self.t = 0

```

```

        self.done = False
        self.profits = 0
        self.positions = []
        self.position_value = 0
        self.history = [0 for _ in range(self.history_t)]
        return [self.position_value] + self.history # obs

def step(self, act):
    reward = 0

    # act = 0: stay, 1: buy, 2: sell
    if act == 1:
        self.positions.append(self.data.iloc[self.t, :]['Close'])
    elif act == 2: # sell
        if len(self.positions) == 0:
            reward = -1
        else:
            profits = 0
            for p in self.positions:
                profits += (self.data.iloc[self.t, :]['Close'] - p)
            reward += profits
            self.profits += profits
            self.positions = []

    # set next time
    self.t += 1
    self.position_value = 0
    for p in self.positions:
        self.position_value += (self.data.iloc[self.t, :]['Close'] - p)
    self.history.pop(0)
    self.history.append(self.data.iloc[self.t, :]['Close'] - self.data.iloc[(self.t-1), :]['Close'])

    # clipping reward
    if reward > 0:
        reward = 1
    elif reward < 0:
        reward = -1

    return [self.position_value] + self.history, reward, self.done

```

```

In [9]: df = pd.read_csv(root_path + "googl.us.txt")
df['Date'] = pd.to_datetime(df['Date'])
data = df.set_index('Date')

```

```

env = Environment1(df)
print(env.reset())
for _ in range(3):
    pact = np.random.randint(3)

```

[illegible]

```
train = df[:delta_split]
test = df[delta_split:]
print("Training date length", len(train), "testing date length", len(test))
```

```
In [12]: import chainer
import chainer.functions as F
import chainer.links as L
import copy
import time
```

```
class Q_Network(chainer.Chain):
```

```
def __call__(self, x):
    h = F.relu(self.fc1(x))
    h = F.relu(self.fc2(h))
    y = self.fc3(h)
    return y
```

```
Q = Q_Network(input_size=env.history_t+1, hidden_size=100, output_size=3)
# GPU Option
#Q.to_gpu()
```

14

```

# Hyper-parameters
epoch_num = 80
step_max = len(env.data)-1
memory_size = 200
batch_size = 20
epsilon = 1.0
epsilon_decrease = 1e-3
epsilon_min = 0.1
start_reduce_epsilon = 200
train_freq = 10
update_q_freq = 20
gamma = 0.97
show_log_freq = 5

memory = []
total_step = 0
total_rewards = []
total_losses = []

start = time.time()
for epoch in range(epoch_num):

    pobs = env.reset()
    step = 0
    done = False
    total_reward = 0
    total_loss = 0

    while not done and step < step_max:

        # select act
        pact = np.random.randint(3)
        if np.random.rand() > epsilon:
            pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
            pact = np.argmax(pact.data)

        # act
        obs, reward, done = env.step(pact)

        # add memory
        memory.append((pobs, pact, reward, obs, done))
        if len(memory) > memory_size:
            memory.pop(0)

        # train or update q
        if len(memory) == memory_size:

```

```

if total_step % train_freq == 0:
    shuffled_memory = np.random.permutation(memory)
    memory_idx = range(len(shuffled_memory))
    for i in memory_idx[:batch_size]:
        batch = np.array(shuffled_memory[i:i+batch_size])
        b_pobs = np.array(batch[:, 0].tolist(), dtype=np.float32).reshape(batch_size)
        b_pact = np.array(batch[:, 1].tolist(), dtype=np.int32)
        b_reward = np.array(batch[:, 2].tolist(), dtype=np.int32)
        b_obs = np.array(batch[:, 3].tolist(), dtype=np.float32).reshape(batch_size)
        b_done = np.array(batch[:, 4].tolist(), dtype=np.bool)

        q = Q(b_pobs)
        maxq = np.max(Q_ast(b_obs).data, axis=1)
        target = copy.deepcopy(q.data)
        for j in range(batch_size):
            target[j, b_pact[j]] = b_reward[j]+gamma*maxq[j]*(not b_done[j])
        Q.reset()
        loss = F.mean_squared_error(q, target)
        total_loss += loss.data
        loss.backward()
        optimizer.update()

if total_step % update_q_freq == 0:
    Q_ast = copy.deepcopy(Q)

# epsilon
if epsilon > epsilon_min and total_step > start_reduce_epsilon:
    epsilon -= epsilon_decrease

# next step
total_reward += reward
pobs = obs
step += 1
total_step += 1

total_rewards.append(total_reward)
total_losses.append(total_loss)

if (epoch+1) % show_log_freq == 0:
    log_reward = sum(total_rewards[((epoch+1)-show_log_freq):])/show_log_freq
    log_loss = sum(total_losses[((epoch+1)-show_log_freq):])/show_log_freq
    elapsed_time = time.time()-start
    print('\t'.join(map(str, [epoch+1, epsilon, total_step, log_reward, log_loss])))
    start = time.time()

return Q, total_losses, total_rewards

```

In [13]: # Let's do it!



```
Q, total_losses, total_rewards = train_dqn(Environment1(train))
```

5	0.0999999999999992	11660	-285.2	51965.05232162122	117.4800
10	0.0999999999999992	23320	-80.8	1027.4448142735287	115.8400
15	0.0999999999999992	34980	-74.4	3630.1065551162233	119.9200
20	0.0999999999999992	46640	-7.6	185.28804305889645	116.4130
25	0.0999999999999992	58300	1.0	94.8101543105673	125.7991000
30	0.0999999999999992	69960	6.0	120.34398262030444	120.1654000
35	0.0999999999999992	81620	12.0	119.88713874976384	122.6150000
40	0.0999999999999992	93280	31.0	67.79507062511402	144.3807000
45	0.0999999999999992	104940	5.8	44.29382843388188	156.4024000
50	0.0999999999999992	116600	-15.6	201.66404752623004	173.0800000
55	0.0999999999999992	128260	-12.8	28.744812113972795	161.4000000
60	0.0999999999999992	139920	-10.8	27.08253281260819	146.2800000
65	0.0999999999999992	151580	-13.0	29.144518221883402	172.4000000
70	0.0999999999999992	163240	-12.0	30.49970243915491	149.4500000
75	0.0999999999999992	174900	-12.2	24.998096223475294	139.9200000
80	0.0999999999999992	186560	-14.0	29.336842975529276	157.9200000

```
In [16]: plt.plot(total_losses,"k",label="loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
```

```
fig = matplotlib.pyplot.gcf()
```


```
fig.set_size_inches(20,10)
plt.legend()
plt.show()
```



```
In [18]: plt.plot(total_rewards,"b",label="Reward")
plt.xlabel("Epoch")
plt.ylabel("Reward")

fig = matplotlib.pyplot.gcf()

fig.set_size_inches(20,10)
plt.legend()
plt.show()
```



```
In [30]: train_env,test_env = Environment1(train), Environment1(test)

# train
pobs = train_env.reset()
train_acts = []
train_rewards = []

for _ in range(len(train_env.data) - 1):
    pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
    pact = np.argmax(pact.data)
    train_acts.append(pact)

    obs, reward, done = train_env.step(pact)
    train_rewards.append(reward)

    pobs = obs

train_profits = train_env.profits
```

```

# test
pobs = test_env.reset()
test_acts = []
test_rewards = []

for _ in range(len(test_env.data) - 1):
    pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
    pact = np.argmax(pact.data)
    test_acts.append(pact)

    obs, reward, done = test_env.step(pact)
    test_rewards.append(reward)

    pobs = obs

test_profits = test_env.profits
pobs = train_env.reset()
train_acts = []
train_rewards = []

for _ in range(len(train_env.data) - 1):
    pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
    pact = np.argmax(pact.data)
    train_acts.append(pact)

    obs, reward, done = train_env.step(pact)
    train_rewards.append(reward)

    pobs = obs

train_profits = train_env.profits

# test
pobs = test_env.reset()
test_acts = []
test_rewards = []

for _ in range(len(test_env.data) - 1):
    pact = Q(np.array(pobs, dtype=np.float32).reshape(1, -1))
    pact = np.argmax(pact.data)
    test_acts.append(pact)

    obs, reward, done = test_env.step(pact)
    test_rewards.append(reward)

    pobs = obs

test_profits = test_env.profits

```

```

In [43]: train_copy = train_env.data.copy()
        test_copy = test_env.data.copy()
        train_copy['act'] = train_acts + [np.nan]
        train_copy['reward'] = train_rewards + [np.nan]
        test_copy['act'] = test_acts + [np.nan]
        test_copy['reward'] = test_rewards + [np.nan]
        train0 = train_copy[train_copy['act'] == 0]
        train1 = train_copy[train_copy['act'] == 1]
        train2 = train_copy[train_copy['act'] == 2]
        test0 = test_copy[test_copy['act'] == 0]
        test1 = test_copy[test_copy['act'] == 1]
        test2 = test_copy[test_copy['act'] == 2]

In [82]: y1 = test0['Close']
        y2 = test["Close"]
        y2 = y2[:len(y1)].values


        plt.plot(np.arange(0,len(y1),1),y1,"r",label="Prediction Q-learning")
        plt.plot(np.arange(0,len(y1),1),y2,"k",label="Data")
        plt.plot(np.arange(0,len(y1),1),y1-y2,"b",label="Difference")
        plt.xlabel("Day")
        plt.ylabel(r"${\rm value}$")
        plt.suptitle("Stock price prediction MAE=%.4f"%np.nanmean(abs(y1-y2[:len(y1)])))

        fig = matplotlib.pyplot.gcf()

        fig.set_size_inches(20,10)
        plt.legend()
        plt.show()

```

Stock price prediction MAE=6.5584



In []:

In []:

In []: